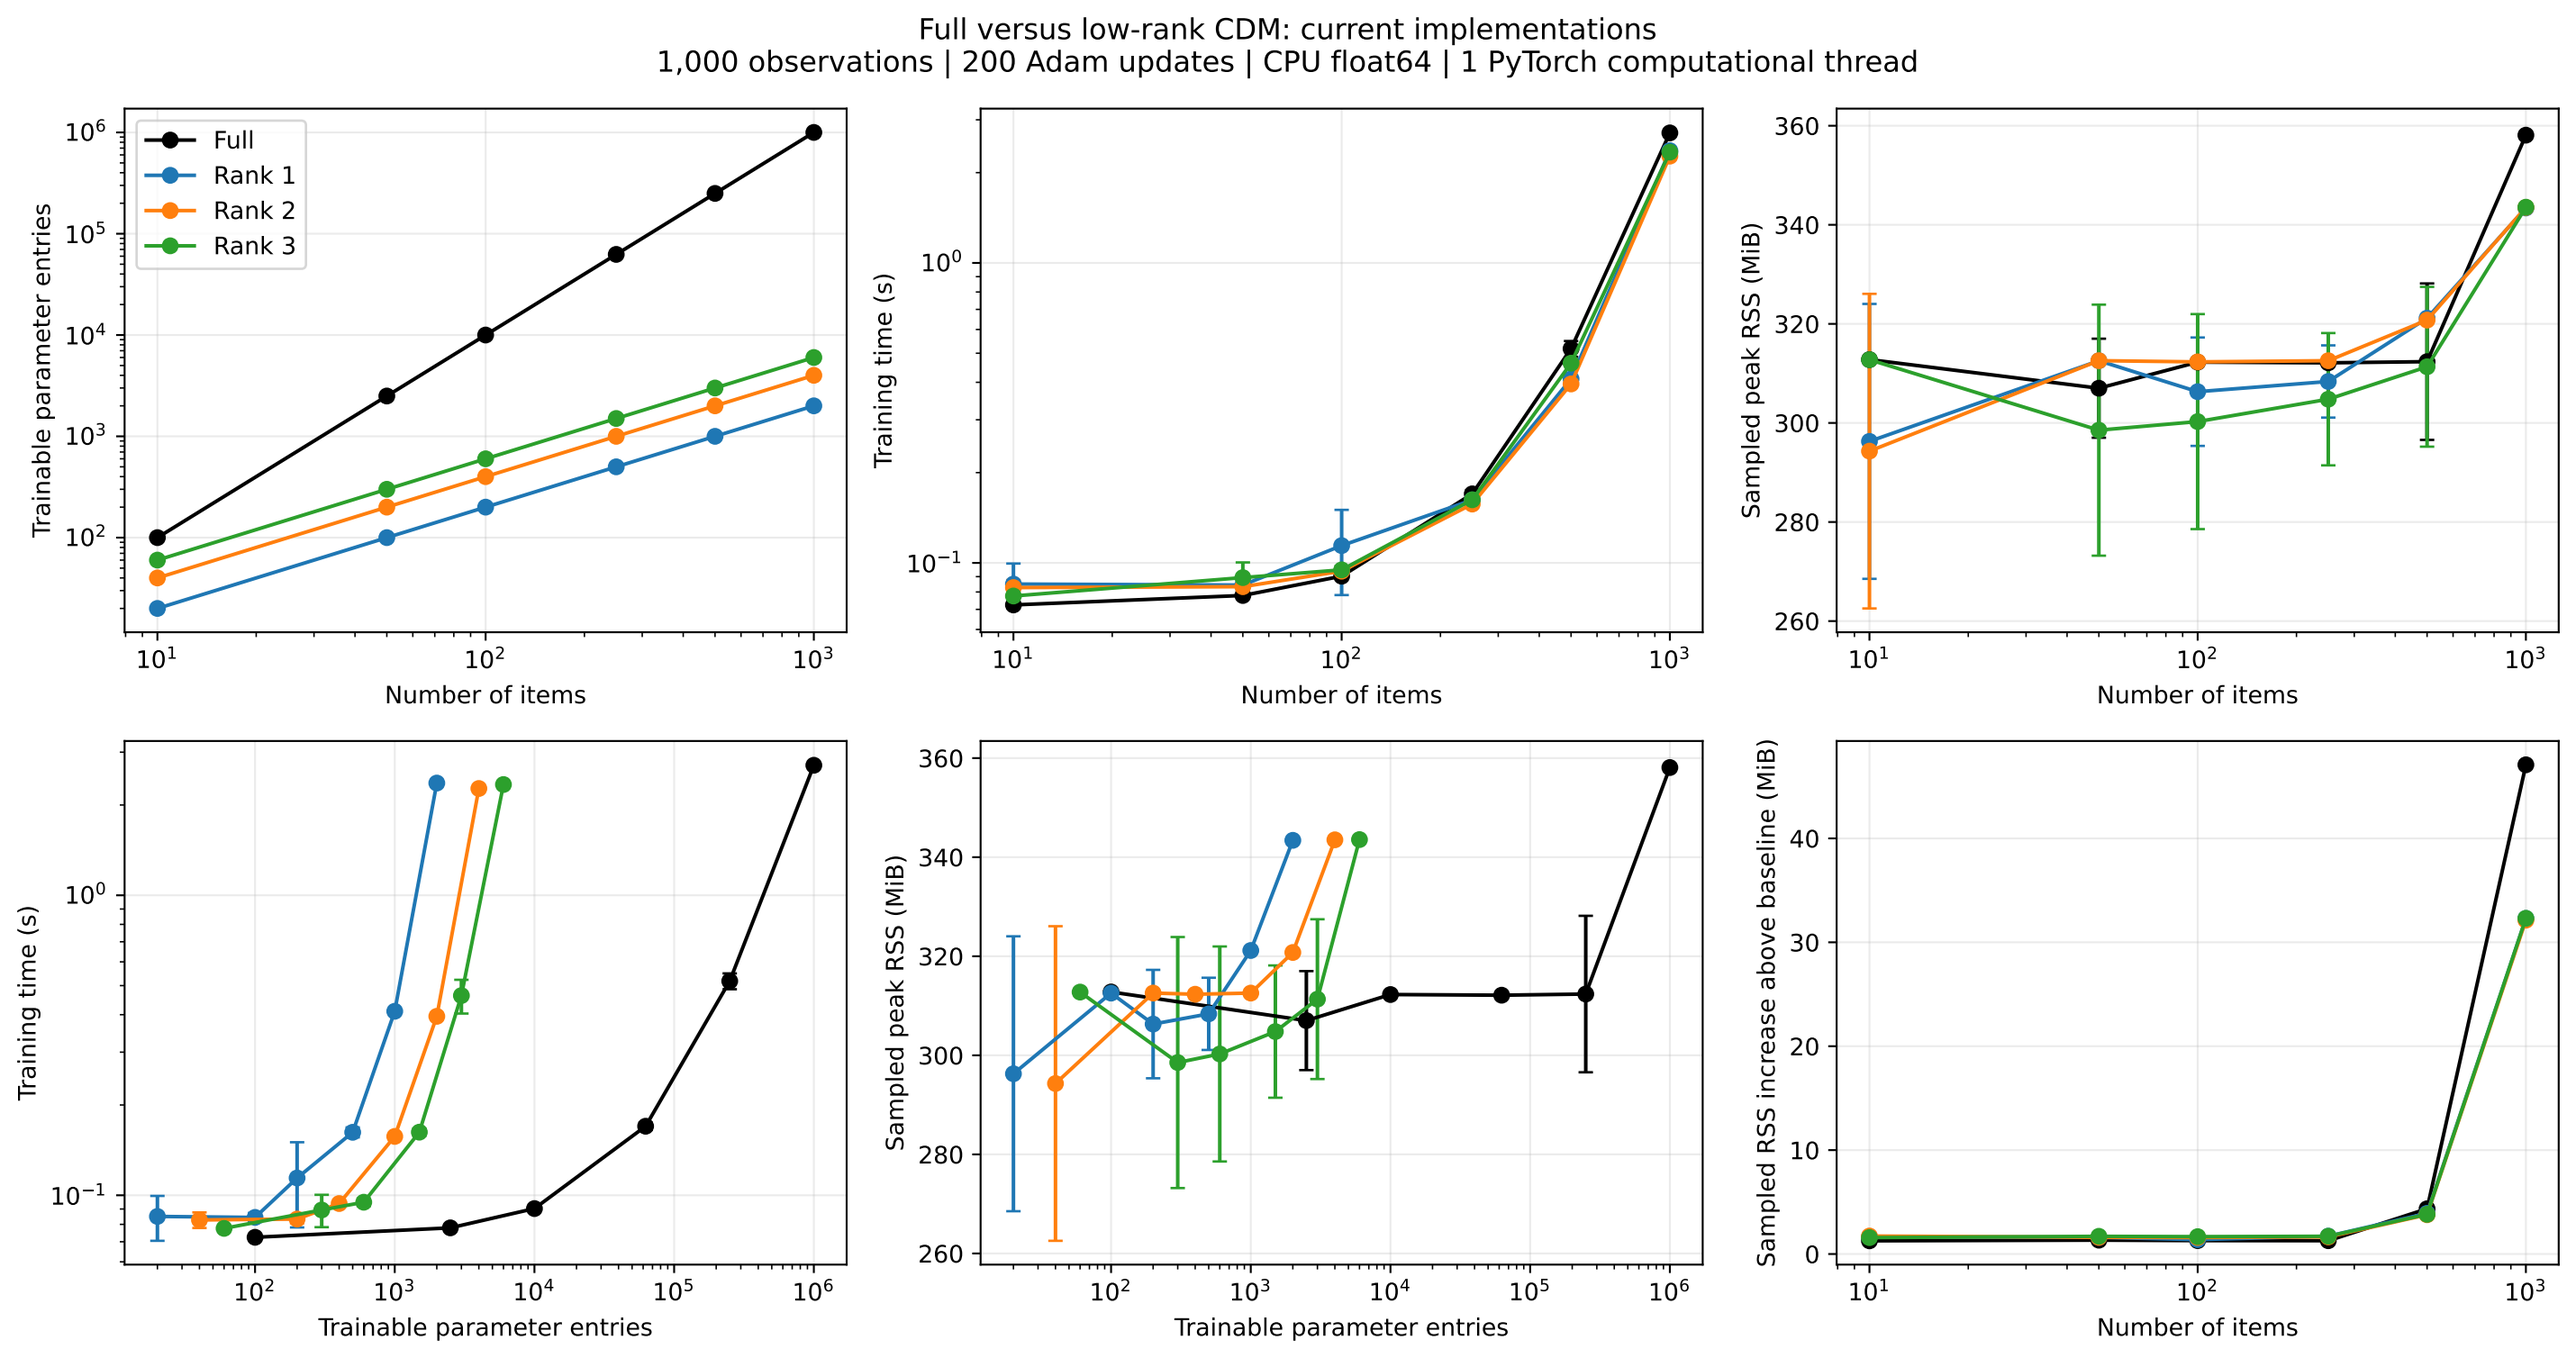
Full versus low-rank CDM: current implementations
1,000 observations | 200 Adam updates | CPU float64 | 1 PyTorch computational thread
Full
Rank 1
Rank 2
Rank 3
Trainable parameter entries
Number of items
Training time (s)
Sampled peak RSS (MiB)
Sampled RSS increase above baseline (MiB)